
PLAYBOOK · \$5,000

IKWE.AI

AI Risk Playbook & Benchmark Pack

Transferable risk intelligence for AI teams

This Playbook provides your team with the complete Ikwe risk evaluation framework, unredacted illustrative audit, sector-specific risk mappings, and reusable benchmark templates. System names are anonymized. Company-specific implementation requires a Full Audit (\$25K).

| | | | |
|--------------------|---|---------------------|--------------------------|
| Contents | Full illustrative audit + sector mappings | Access Level | Playbook |
| Methodology | EQ Safety Benchmark v2.1 | Transferable | Yes — any team can apply |

This Playbook includes the language, structure, and evidence. It does not include your audit.

ikwe.ai/audit · stephanie@ikwe.ai

PLAYBOOK CONTENTS

What's in This Pack

■ Full Illustrative Audit (Anonymized)

Complete 9-page audit with all sections visible — findings, scorecard, risk events, failure modes, remediation framework, priority actions. System names anonymized.

■ Sector-Specific Risk Mappings

Risk dimension weightings for Healthcare, Policy/Gov, LLM/Agents, Workplace/HR, and Consumer/Platform sectors.

■ EQ Safety Benchmark Framework

Methodology documentation: Safety Gate layer, 8 weighted quality dimensions, scoring rubric, severity/liability/scale formula.

■ Reusable Scorecard Templates

Blank risk scorecard, risk event log, failure mode map, and audit evidence templates. Ready to populate for your systems.

■ Before/After Remediation Playbook

The exact fix-by-fix framework showing how each risk dimension was reduced 42-67%.

■ YOUR Company-Specific Audit

Requires Full Audit (\$25,000) — we evaluate your systems, your data, your risk.

■ Board-Ready Report (Your Company)

Requires Full Audit (\$25,000) — named, scored, defensible for your board.

■ Ongoing Governance & Re-Audits

Requires Implementation (\$50K+) — we stay on top of it so you don't have to.

This gives you the language and structure. It does not give you your audit.

How Risk Weights Shift by Sector

Different sectors create different risk profiles. The same AI behavior carries different weight depending on context, regulatory environment, and user vulnerability.

| Risk Dimension | Healthcare | Policy/Gov | LLM/Agents | Workplace | Consumer |
|----------------------|------------|------------|------------|-----------|----------|
| Emotional Escalation | ■ HIGH | ■ LOW | ■ LOW | ■ MED | ■ HIGH |
| Dependency Formation | ■ MED | ■ LOW | ■ MED | ■ MED | ■ HIGH |
| Authority Drift | ■ HIGH | ■ HIGH | ■ MED | ■ HIGH | ■ MED |
| Scale Amplification | ■ MED | ■ MED | ■ HIGH | ■ LOW | ■ HIGH |
| Governance Failure | ■ MED | ■ HIGH | ■ HIGH | ■ MED | ■ MED |

■ HIGH = primary risk weight · ■ MED = secondary · ■ LOW = monitor

Healthcare AI

Therapy-adjacent, crisis support, emotional mirroring systems

Primary risks: Emotional Escalation + Authority Drift

Policy & Government

Decision-support, citizen-facing, regulatory-adjacent systems

Primary risks: Authority Drift + Governance Failure

LLM & Agents

Autonomous agents, tool-using AI, multi-step reasoning systems

Primary risks: Scale Amplification + Governance Failure

Workplace & HR

Employee assistants, performance tools, internal copilots

Primary risks: Authority Drift + Dependency Formation

Consumer & Platform

Public-facing chatbots, social AI, entertainment systems

Primary risks: Scale Amplification + Emotional Escalation

Note: The Full Audit (\$25,000) applies these sector weightings to YOUR specific systems and generates a company-specific risk report.

Ready for your company-specific audit?

Full Audit: \$25,000 · 4-week delivery · Board-ready report

ikwe.ai/audit