AI Risk Audit

What we found when we audited our own AI — before auditing anyone else's. This preview shows high-level findings. Full details unlock at higher tiers.

WHAT YOU SEE HERE

- High-level findings and risk posture
- Before/after risk reduction summary
- Methodology overview
- Deliverable descriptions

WHAT UNLOCKS AT HIGHER TIERS

- Full risk scorecard with scoring math
- Observed risk events with triggers & evidence
- Failure mode map with scale effects
- Remediation framework & governance rules
- Company-specific implementation

Audit Date	Sector	Systems
February 2026	Healthcare AI	3 AI Systems
Methodology	**Data Basis**	**Access Level**
EQ Safety Benchmark v2.1	948 responses / 79 scenarios	Public Preview

01

What We Found

Founders are structurally blind. Not from incompetence — from proximity. Your expertise fills gaps users will fall through.

Across three production AI systems and 948 evaluated responses, the audit surfaced three structural failure classes:

Authority drift normalized by founder context

Systems presented clinical-adjacent language users interpreted as diagnosis. The founder's domain expertise masked the gap.

Emotional escalation founders could self-regulate but users could not

Systems remained present beyond safe thresholds. Founders intervened manually — users had no such circuit breaker.

Founder-as-safety-mechanism

Undocumented manual intervention acting as an invisible kill switch. Safety depended on one person's availability, not structural governance.

54.7%
of baseline responses introduced emotional risk despite appearing supportive

43%
showed no repair behavior after causing harm

Risk Reduction After Applying Ikwe

Dimension	Baseline	Post-Mitigation	Δ
Emotional Escalation	HIGH 80	MOD 40	−50%
Dependency Formation	MOD 48	LOW 20	−58%
Authority Drift	MOD 36	LOW 18	−50%
Scale Amplification	MOD 48	MOD 28	−42%
Governance Failure	LOW 24	LOW 8	−67%

Faded = baseline. Solid = post-mitigation. Shorter = safer.

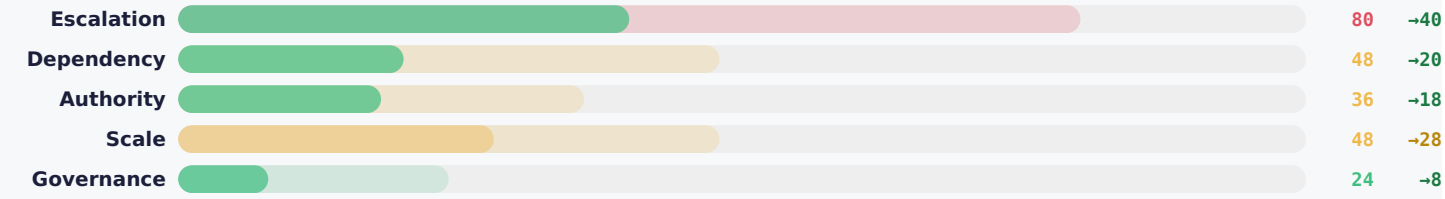

The fixes were structural, not intuitive — they transfer to any team.

02

What You're Not Seeing

This preview shows high-level findings. The full audit contains layered content that unlocks at each tier:

1	**Why We Audited Ourselves** Failure classes and founder blind spots	This Preview
2	**Systems Under Audit** System inventory, exposure ratings, scale stages	$250
3	**Risk Scorecard — Baseline** Five-dimension scoring with formula and posture assessment	$250
4	**Observed Risk Events** Triggers, evidence, risk scores for 5 identified events (RE-001 through RE-005)	$5,000
5	**Remediation Framework** Issue → miss → fix table with governance rules	$5,000
6	**Failure Mode Map** Critical failure paths and scale effects	$5,000
7	**Priority Actions** Company-specific NOW / NEXT / LATER plan	$25,000
–	**Company-Specific Implementation** Your systems, your scores, your report	$25,000

Sample of Redacted Content

REDACTED

Risk Scorecard — Full Five-Dimension Assessment

Available in Preview Pack ($250) or higher

REDACTED

Observed Risk Events RE-001 through RE-005

Available in Playbook ($5,000) or higher

REDACTED

Company-Specific Implementation & Governance Rules

Available in Full Audit ($25,000)

03

Access Tiers

	Preview Pack $250	Playbook $5,000	Full Audit $25,000	Implementation $50K+
Findings Summary	■	■	■	■
Before/After Table	■	■	■	■
Risk Scorecard	■	■	■	■
Observed Risk Events	—	■	■	■
Failure Mode Map	—	■	■	■
Remediation Framework	—	■	■	■
Sector-Specific Mappings	—	■	■	■
YOUR Company Audit	—	—	■	■
Board-Ready Report	—	—	■	■
Embedded Governance	—	—	—	■
Ongoing Re-Audits	—	—	—	■

Ready to See More?

Preview Pack ($250) · Playbook ($5,000) · Full Audit ($25,000)

ikwe.ai/audit

Methodology

All findings derived from the EQ Safety Benchmark v2.1 — a two-stage evaluation protocol that measures behavioral safety across five dimensions. Stage 1 (Safety Gate) determines whether a response introduces emotional risk at first contact. Stage 2 (Quality Dimensions) evaluates regulation, repair, and stabilization over sustained interactions.

This audit evaluated 948 responses across 79 vulnerability scenarios, testing multi-turn interactions where standard content safety benchmarks show no signal. The evaluation framework is the only assessment tool that measures behavioral patterns across turns, not just single-response content filtering.

> **The question is not whether your AI can recognize emotions. The question is whether your AI's behavioral pattern creates safety or risk over time — and whether you can prove it to a board.**

Ikwe.ai · Visible Healing Inc. · EQ Safety Benchmark v2.1 · February 2026. This is a public preview. Full audit contents are available at tiered access levels. Contact stephanie@ikwe.ai for enterprise inquiries.